

Kavya Kushnoor

Work Eligibility: U.S. Citizen

📍 Portland, OR

✉ k.kavya101@gmail.com

📞 804.484.0170

EDUCATION

UNIVERSITY OF NORTH CAROLINA, WILMINGTON

MS, COMPUTER SCIENCE Dec, 2019

VIRGINIA COMMONWEALTH UNIVERSITY

MBA, FINANCE Aug, 2015

MANIPAL UNIVERSITY, INDIA

BBM, E-COMMERCE Dec, 2011

LINKS

<https://kavyakushnoor.github.io>

<https://linkedin.com/in/kavya-kushnoor>

www.github.com/kavyakushnoor

CERTIFICATIONS

- AWS Cloud Practitioner
- Containers with Docker, K8s & OpenShift
- Django App Development with SQL & Databases
- Machine Learning with AWS
- Machine Learning for Trading
- Developing AI apps with Python & Flask
- Python for Data Science, AI & Development

SKILLS

PROGRAMMING:

Python, JavaScript, Java, Go, C++, Bash

CLOUD SERVICES:

AWS, Microsoft Azure

MACHINE LEARNING:

AWS Sagemaker, XGBoost, Pytorch, Sklearn

FRAMEWORKS:

Django, Flask, AWS Lambda

DATA:

SQL, DynamoDB, Postgres, NumPy, Pandas

CI/CD:

Github Actions, Circle CI, Terraform, Docker, K8s

EXPERIENCE

Cloud Solutions Engineer

Intel, Hillsboro, OR

May 2021 - Feb 2023

- Led the Postgres benchmark for a software defined storage solution against direct-attached SSDs; planned architecture and hardware set-up; collected performance and telemetry data, did performance analysis, published technical blog for marketing and product visibility
- Led the framework development of full-stack API integrated web application to gauge customers' software ecosystem and reflect market research using Python and JavaScript; improved efficiency and reduced human error by automating reporting processes up to ~80%; mentored team member by integrating them into project
- Automated performance testing programs and benchmarking by integrating telemetry tools into benchmark scripts using Python, Go, Java and Bash to fulfill CPU quality requirements; Developed, debugged and validated performance tools such as health-inspector and perkitbenchmarker along with Docker and Kubernetes; documented instructions/tutorials regarding tool usage; Redfish REST API based schemas to gauge server health.
- Identified workload bottlenecks and applied software optimizations to increase workload efficiency by 1.8x-4.6x, such as version and kernel upgrades, applying software accelerators (e.g., QAT) and patches, recommending HW upgrades by examining server health via performance tools, implementing jemalloc, huge pages setup, etc.
- Supported customer requests to help identify suitable database proxy workload matching the db performance in customer's production environment. Reported throughput and CPU performance data to enable workload comparison using automated scripts based on live customer feedback and requirements; gauged the benchmarking app with compute-based configurations and storage optimized solutions recommended Linux and windows set-up, tuning and functional use; shell scripting

Cloud Solutions Engineering Intern

Intel, Hillsboro, OR

May 2019 - Aug 2019

- Manually tested and reported scripts generated by the server platforms at the Data Center Group in Linux
- Developed Python scripts to automate Rack Scale Design (RSD) Conformance Testing for Redfish REST API based schemas to gauge real-time health of RSD setup in labs
- Worked with backend engineering team for customer requirement, product improvement and bug fixes
- Revised, edited and upgraded existing documentation based on test results
- Collaborated with teammates across several time zones and geographic locations for framework development

Software Developer

Center for Innovation & Entrepreneurship, Wilmington, NC

Aug 2016 - Dec 2019

- Developed website content as a software service for education and writing using HTML, CSS, JavaScript, Java and bootstrap; research project in UX, Human-computer interaction and design. Gathered requirements and updated website content regularly.
- Developed responsive application using Alexa. Voice Service for end-user speech practice that interacts with users in real time using manually created objects with AWS Lambda, AWS DynamoDB and API Gateway which controls the display of data on a web browser.
- Worked on signal processing, image processing and steganography-based assignments.

Digital Technology Leadership Program - Intern

General Electric, Glen Allen, VA

May 2017 - Aug 2017

- Developed application using python and bash scripts to automate deployment of security-guard app on multiple PCs in the organization; wrote a Python script that automates the process of pinging user systems, updating security app version and clearing previous version from the system based on the users' IP address
- Created Splunk dashboards to visualize the real-time health of the power equipment and performance; met with internal customers and gathered feedback about data needs; assimilated feedback into the final reports
- Facilitated automated delivery of reports to internal customers; developed, debugged and validated applications based on test results and feedback